

ProjectReport

November 21, 2020

1 Abstract

Our project goal was to build a functional movie recommendation system using publicly available data. To this end, we used the [GroupLens MovieLens 25M dataset](#), and built a web scraper to extract other information from IMDb. We cleaned, filtered and analysed our data to come up with a workable plan for constructing an adjacency matrix of movies. This is constructed by taking a linear combination of various similarity matrices made using different similarity parameters. The coefficients of the linear combination were calculated by running gradient descent multiple times to minimize an error function. We used anonymised user ratings data to subset this matrix based on user input, run PageRank on the subset, and return the highest recommended movies. We also built a GUI for this application.

2 Introduction

Recommender systems are algorithms that attempts to predict user preferences. For example, recommender systems are used widely by services such as Spotify and Netflix to recommend songs and movies respectively to users, based on their past preferences in the respective domains. Recommender systems have grown significantly, and they are core parts of various online content providers to the extent that said content providers offer significant rewards for improvements to their algorithms. This can be seen in [Netflix Prize](#).

PageRank is an algorithm used to rank importance of nodes in a network. It works on the concepts of random walks, wherein a fictional person traverses the graph at random, and keeps track of the frequency of all the nodes visited. This eventually converges, and the frequency values represent relative importance of the nodes in the graph.

2.1 Problem

The aim for this project was to build a PageRank based movie recommendation system from first principles. By defining a similarity parametric for all pairs of movies using publicly available data, and thus creating an adjacency matrix, we can find the most important movies in a network using PageRank. Methods of making the recommendation user-specific are discussed later.

2.2 Literature

Our main inspiration and source of information were interactive sessions of CS522 course held during this semester. From there we learnt about PageRank and various methods of computing it.

2.3 New Idea

After learning about PageRank, and its implementations, we attempted to use the power iteration method of calculating PageRank to attempt to build a movie recommendation system.

3 Method

3.1 Implementation Details

3.1.1 Data Collection and Processing

The data for this project is all from publicly available sources. Our first goal was getting a comprehensive set of movies. Initially, we tried getting [IMDb data](#), but realised this is only a very small and scattered subset. Eventually, we found and decided to use the [GroupLens MovieLens 25M dataset](#), present in the `ml-25/` directory. This consists of a relatively small set of around 62k movies, but this was more than enough for our purposes. In addition to just movies, this dataset also included for each movie its title, IMDb ID, year of release, genre(s), tags, and perhaps most importantly 25 million (hence the name) anonymised user ratings for the movies. For additional details about the specific data contained in this dataset, refer to `ml-25m/README.txt`.

For additional data about each movie, we built a web scraper that used the IMDb ID to go to the IMDb URL and extract the relevant information. Information such as **Genre**, **Relevant Tags**, **Director**, **Cast**, **Language** and **Rating** was scraped from the IMDb website for each movie. `Requests` was used to get the webpage in HTML format, then `BeautifulSoup4` was used to scrape the required data. The script for the scraper can be found in `scraped-movies-data/scrapper/web_scraper.py`.

To clean the movie data and collect it into a proper format, `pandas` and `re` libraries were used to handle the csv files and process text patterns respectively. This process is present in `cleaned_movies_generator.py`, and its output is present in `ml-25m/clean_movies.csv`. Additionally, most of the movies in the final dataset were irrelevant, primarily because they were of other languages. Additionally, we could not process the entirety of the data. For example, a single 62000×6200 matrix of half-precision (2 byte) floats that would represent similarity between movies would take up over 7GB of space. To this end, movies were filtered to include only those in English, Hindi, Urdu and Punjabi. Additionally, there were some movies whose IMDb IDs were invalid (the webpage resulted in a 404 error while scraping) and these movies were dropped. This resulted in a set of 23843 movies, which are available in `processed_data/cleaned_subsetted_movies.csv`.

Additionally, `ml-25m/ratings.csv` contained in addition to `userID`, `movieID`, and `rating`, the `timestamp` for each rating. Since this was of no use to us, the column was dropped and the resulting dataset is in `ml-25m/timeless-ratings.csv`.

3.1.2 Methodologies Employed

Graph-centric Initially, we tried a graph-centric approach. In this method, the movies would be nodes on a weighted, undirected graph, and edges would indicate similar movies, the edge weight being the relative similarity for each movie. We chose `graph-tool` as our library of choice since in comparison to various other libraries, it came out significantly faster in various [Benchmark tests](#).

The plan was to use various parameters of the collected data to create graphs that represent those similarities. For example, using the user ratings a graph was created

(`graphs/ratings_graph.gt.xz`) in which users and movies were linked by edges weighted with the rating given to that movie by that user. This was intended to be used to calculate one similarity parametric between two movies, by using the fact that similar movies would have fewer edges connecting them, and these edges would have higher weight. This allowed us to then create a parametric, where the similarity between two movies ($S_{u,v}$) is directly proportional to the number of shortest paths ($N_{u,v}$) (in terms of number of edges) between two movies, directly proportional to the average weight along the paths ($A_{u,v}$), and inversely proportional to the number of edges in the paths ($E_{u,v}$). Mathematically,

$$S_{u,v} = \frac{N_{u,v} A_{u,v}}{E_{u,v}}$$

This approach quickly fell apart when we constructed the graph, which turned out to be extremely slow to process, and impossible to run many centrality or relatedness algorithms on simply due to the large memory requirements.

Matrix-centric After the lecture on PageRank and calculating it using power iteration, we attempted to represent movie similarities as matrices. For this, we identified parameters using which we can compare movies. These are director, cast, tags, genre, user ratings and language. We attempted to use `h5py` library that allows for storing large matrices on disk as `.hdf5` files. While this method would have worked in concept, the processing times were too long since reading from and writing to disk is a very slow process. It was at this point when we realised that 62000 is too large a set of movies to compute on, even with low-precision datatypes. As a result, we subsetting the movies as was described above.

Our goal at this point was to calculate one matrix for each similarity parameter, and take a linear combination of these matrices to obtain the final adjacency matrix representing similarities between all pairs of movies. Director, genre, language and rating similarities were quick to calculate. In the case of genres, this was due to the fact that there were very few total genres in the dataset. For director, language, and rating, the process of computing similarities were not computationally expensive.

Cast similarity and user similarity had the same issue, that the number of different users/cast members was too large to compute. Theoretically, our method would benefit from being able to compute the cast similarity as well, however due to restriction of our systems this was not possible. User similarity was incorporated in two different ways. Firstly, the 20,000 users with the most ratings were used to construct a user similarity matrix, which was then used to compute an error function that enabled calculating coefficients of the linear combination through gradient descent. This method is explained in more detail in the next section. The second application is used to subset the movies to run pagerank on. This method is explained in a later section.

All the similarity matrices had different ranges of values, depending on how they were calculated. Except for director and language similarities, the matrices also had the issue that larger values represented greater differences, instead of greater similarities. To rectify this, and normalize the values, 1 was added to each cell of the matrix, and then reciprocal of all values was taken. This normalises the values, since 0 cells along the diagonal and other places (indicating 100% similarity in that metric) become 1, and larger difference values are less than 1. This added a constraint on the linear combination coefficients, that they had to sum to 1 for the scale to remain the same. This is a workable restriction.

All the similarity matrices can be found in the `all_similarities/` directory. The code to calculate rating similarity is in `all_similarities/calculation/rating-similarity.py`. For director similarity, it is in `director-similarity.py` in the same directory. The same process was applied for language similarity. Genre and tags similarity was calculated using the method in `similarity-calculator.py`. User similarity was calculated using `user-similarity-calculator.py`. Raw, testing code can be found in `movie-similarity-generation.ipynb`. The postprocessing code to add 1 to and calculate reciprocal of the matrices is also present there.

3.1.3 Calculation of Linear Combination Coefficients

To calculate the coefficients for the linear combination coefficients of the similarity matrices, gradient descent was employed. The user similarity matrix calculated previously was used as a target matrix. The error function was calculated as the sum of absolute difference between corresponding values of the adjacency matrix and the user matrix. Using `scipy.optimize.minimize` function and `method='BFGS'` gradient descent was run from multiple different starting values, and the best set of weights obtained was used for the final adjacency matrix. The code for this is available at `weights-gradient-descent.py`.

3.1.4 Final Recommendation System

After the adjacency matrix was obtained from the coefficients calculated through gradient descent, PageRank was implemented using the power iteration method. The input movie names are matched to the ones in our dataset through fuzzy matching, using `fuzzywuzzy` library for Python. This operates by calculating the [Levenshtein Distance](#) for each title in our database, given an input title, and picks the one with the highest value.

After matching movie titles, these are mapped to their respective `movieId` values. This is given as input to `user_id.subset()` which...

!! TODO Harsh, explain about user-id

A collection is then formed comprising of all the movies rated by all the given users. The users are returned such that the number of movies is close to 3000. These `movieId` values are given to the `pagerank` function, which subsets the adjacency matrix accordingly and performs PageRank using the power iteration method.

The starting vector for power iteration is taken as $v = (1, 0, 0, \dots, 0)$, and convergence is achieved when the maximum difference between corresponding values in successive iterations of v is less than a configurable tolerance ϵ . Mathematically, convergence occurs if

$$\max(v_i - v_{i-1}) \leq \epsilon$$

Where v_i is the value of the vector v after the i th iteration. A cap on iterations was implemented as well, to prevent an infinite loop in the rare case that convergence does not occur.

The function returns a dictionary mapping `movieId` to PageRank, of which the 5 largest values are the recommended movies for the given input.

3.1.5 Tools and Libraries Used

All the code for this project is written in Python 3.8.5. The libraries used in this project and their usages are as follows

Library	Usage	Link
numpy==1.19.4	Various mathematical applications, including handling of very large matrices, matrix multiplication, and various other matrix operations	https://numpy.org/
scipy==1.3.3	Used <code>scipy.optimize</code> to perform gradient descent and find optimal combination of weights for different similarity parameters	https://www.scipy.org/
pandas==1.1.3	For handling and processing of raw and processed data, in the form of <code>.csv</code> files	https://pandas.pydata.org/
re	Python Regular Expressions library for pattern matching while processing raw data	https://docs.python.org/3/library/re.html
numba==0.51.2	A Python compiler that allows for compiling a subset of python and numpy to machine code, used while testing methods and processing data to accelerate some computations	https://numba.pydata.org/
graph-tool	A Python library written in C++ for handling graphs, used while testing different methods of PageRank and methods of processing movie similarity. This was eventually not used in the final result	https://graph-tool.skewed.de/
bs4	Web scraping	link
lxml	Web scraping	link
fuzzywuzzy==0.18.0	Used for fuzzy matching of input movie names to those in our dataset	https://pypi.org/project/fuzzywuzzy/
python-Levenshtein==0.12.0	Used by <code>fuzzywuzzy</code> to accelerate calculation of Levenshtein distance	https://github.com/ztane/python-Levenshtein/

4 Results

4.1 Experiment Findings

TODO : UPDATED WEIGHTS, EXPLANATION

The final adjacency matrix obtained is `final_adj_mat.npz`. The weights used to obtain this were (0.873, 0.047, 0.08) for director, genre, and tags similarity respectively. Additionally. It was found that Drama and Comedy movies seem to dominate the recommendations. The recommendation system can be interacted through the GUI app, by running *TODO : INSERT FILE NAME HERE*.

4.2 Interpretation of Findings

Observing the weight values, it appears as though director similarity is given a surprisingly high weightage. This may make some sense, since directors often make similar types of movies. Genres was also surprisingly given a fairly low weightage. This could possibly be because tags serve a similar purpose, and since there are a larger number of tags, the granularity of similarity values would be much more, and consequently less polarising.

Upon exploration, it appeared that Drama and Comedy were very frequent genres in our dataset. Due to this, while using `user-id.subset` they have disproportionately higher influence on the subsetted users. This tends to bias the recommendations towards movies with these genres, since they are the ones subsetted. Currently, movies without these tags have better recommendations.

5 Conclusion

PageRank appears to be a workable technique for building a recommendation system. It may be possible to improve the system further if larger data can be processed, but this project can serve as a proof of concept for such an approach. There are definitely various improvements that can be made to the approaches we have taken to process and use our data, but the current approach has also provided us with workable results. Some improvements could be in the form of finding an appropriate normalization for different frequencies of genres, to avoid biasing. Additionally, the algorithm could be significantly improved if additional data is able to be processed and used in the similarity metric. We did not have a deep understanding of the various methods of gradient descent, and perhaps a better method, implemented appropriately, could have offered better results. In relation to this, there could be significant improvements to the error function used for gradient descent.

This project was a major learning experience, in various aspects of data analysis and recommendation systems. We learnt to scrape web data, process data with pandas and numpy, storing and manipulating large data, build GUIs, and applications of scientific techniques like gradient descent.

5.1 Team Work

This project was definitely a huge group effort. We managed to effectively split the workload by distributing different tasks, and combining the results. Another method of collaboration was splitting the computation. For time consuming processes, like scraping web data, we were able to effectively halve the total time by each doing half the computation. This was applied not only for scraping web data, but also calculating tags and user similarity matrices, and running gradient descent for various input weights to find the best combination of coefficients.

Not only just computation and effort, but ideas are also much better with groups. We were able to effectively find faults in each other's ideas, to identify areas of improvement, and effective methods of building a better recommendation system. The project involved not only a lot of coding, but a

lot of discussions held online to find solutions to various problems encountered. We also spent large amounts of time spent exploring and getting to understand the data, then sharing our findings to think of better pipelines to process our data.